

unidb — Engine Architecture & Product Guide

One embedded database engine that unifies relational tables, vector search, graph relationships, and a durable event stream over a single storage core — so one transaction can save a row, its embedding, a graph edge, and an event **atomically**, in one commit.

The one idea to remember. A typical AI-native app stitches together Postgres + a vector store + a graph database + a message queue, and writes to all four on every change with no shared transaction. unidb collapses that into **one engine and one commit**. Measured against exactly that four-system stack at matched flush-to-disk durability, unidb is **3.61x faster** and loses **zero** data on a crash where the four-system stack tears a record.

Contents

1. What unidb is
2. Architecture at a glance
 1. The layered engine
 2. Core guarantees
 3. Ways to run it
3. How your data is stored & kept safe
 1. On-disk format
 2. How a write becomes permanent
 3. Crash recovery — when the power goes out
 4. Protection against corruption
 5. Memory & space management
4. Transactions & working with many users at once
 1. Snapshots & isolation levels
 2. Reading and writing at the same time
 3. What happens on rollback
5. The SQL layer
 1. Supported data types
 2. Writing queries — examples
 3. Query capabilities
 4. Constraints, security & roles
6. Search, indexing & relationships
 1. Automatic durable indexes
 2. Vector similarity search
 3. Full-text search
 4. Graph relationships
7. Real-time events & streaming
8. The REST server — full API reference
 1. Getting started
 2. Running SQL
 3. Transactions over HTTP
 4. Rows, graph & indexes
 5. Events & realtime
 6. Admin, stats, logs & metrics
 7. Replication
 8. Error codes
9. Operations & high availability
10. Performance — measured against Postgres
11. Correctness & testing
12. Known limitations, bugs & roadmap priority
 1. Known bugs
 2. Known limitations & feature gaps
13. Where unidb is headed

1. What unidb is

unidb is a single embedded database engine written in Rust. Like SQLite, it runs **inside your application** as a library — no separate server process required — and keeps all of your data in one file. Unlike SQLite, it natively handles four kinds of data at once, over one shared storage and transaction core:

- **Relational tables** — a practical SQL dialect over versioned tables: joins, aggregates, subqueries, CTEs, a cost-based query optimizer, and `EXPLAIN [ANALYZE]`.
- **Vector search** — `VECTOR(n)` columns with a durable on-disk similarity index and a `NEAR` operator for nearest-neighbor queries (the building block of semantic / AI search).
- **Graph relationships** — directed edges between entities, with a Cypher-style read query for traversals.
- **Durable event stream** — a built-in, Kafka-style event queue with consumer offsets and replay, for change-data feeds and realtime.

Because all four live in **one engine, one log, and one transaction manager**, a single transaction can touch all of them and either commit everything or nothing. That is the capability an application otherwise has to fake with distributed writes across separate systems — and can never make truly atomic.




Figure 1 — One atomic multi-model commit vs. the four-system stack it replaces.

What unidb is not. It is honest about its scope. It is not trying to beat Postgres on every single-table workload, it does not implement the entire ANSI SQL standard, and it runs as a single primary with read replicas (no distributed multi-writer consensus). Its advantage is the unified, multi-model, single-commit workload — the exact pattern AI-native apps are built on.

2. Architecture at a glance

2.1 The layered engine

Every capability is built as a thin layer on top of the same storage core. That is what makes cross-model atomicity *structural* rather than a promise: a graph edge and a queue event are, underneath, just ordinary rows in the same store, protected by the same log.



Figure 2 — The layered engine. Everything sits on one shared storage core.

2.2 Core guarantees

- **Fully durable (ACID).** Every committed transaction is on disk before the commit returns. A committed change survives a crash; an in-flight one leaves no trace.
- **Runs in your process, thread-safe.** The engine is shared safely across threads; many readers and writers work concurrently and scale with CPU cores. The default embedded build pulls in **no async runtime and no network stack** — it is a plain, synchronous library.
- **Log-before-data, always.** No change ever reaches the data file before it is safely recorded in the write-ahead log. This single rule is what makes crash recovery provably correct.
- **Explicit transactions.** Every operation runs inside a transaction you begin and then commit or roll back — no surprising implicit commits.

2.3 Ways to run it




Figure 3 — Deployment modes: embedded library (primary), optional REST server, Rust attach client, plus single-primary + read-replica high availability.

3. How your data is stored & kept safe

3.1 On-disk format

- **One data file** made of fixed-size **8 KiB pages**, plus a separate, segmented write-ahead log directory. Keeping data in one file keeps backup, copy, and recovery simple.
- **Every page is checksummed** and carries a version stamp, so corruption is detected on read rather than silently returned.
- **Rows are stored in slotted pages.** Each row carries a small header that records which transaction created it and which (if any) superseded it — the basis of unidb's versioning (Section 4). New column types are added without ever changing the on-disk format.




Figure 4 — Page and row layout. Every page is checksummed; every row records its version stamps.

3.2 How a write becomes permanent

unidb uses a write-ahead log (WAL): before a change is applied to a page, a record describing it is appended to the log. The log is the source of truth. A commit is not acknowledged until its log record is physically flushed to disk.

Within a transaction, individual statements append to the log *without* flushing; the single flush happens at `COMMIT`. When many transactions commit at once, the engine **coalesces them behind one shared flush** ("group commit") — so throughput scales with concurrency instead of being bottlenecked by one flush per statement. Read-only transactions skip the flush entirely.

Write path: statement → log → page → group-commit flush




Figure 5 — The write path. Log-before-page is enforced at every change; durability is one shared flush per commit.

3.3 Crash recovery — when the power goes out


If the process is killed or the machine loses power mid-write, you do nothing special: the next time the database is opened, unidb **repairs itself automatically**. Recovery is deterministic and needs no manual intervention.

On open, the engine:

1. Reads the **control file** — a tiny, checksummed record that always points at the last known-good checkpoint. Recovery always starts here.
2. **Replays** committed changes from the log, forward from that checkpoint, so every acknowledged commit is restored.
3. **Rolls back** anything belonging to transactions that never reached a commit, leaving no partial rows behind.

The guarantee. After recovery, the database contains **exactly the transactions that had committed** — no more, no less. No half-applied statements, no orphaned rows, no readable corruption. This holds even if the crash happens *during* recovery itself: re-running recovery is safe and produces the same result.

Automatic recovery when the database re-opens



Even total loss of the data file can be rebuilt from the log alone; a checkpoint periodically trims the log so recovery stays fast.

Figure 6 — Crash recovery on open: control file → replay committed → roll back uncommitted.

3.4 Protection against corruption

- **Torn-page protection.** An 8 KiB page write is not atomic at the hardware level — a crash mid-write can leave a page half-old and half-new. The first time a page changes after each checkpoint, unidb saves a full, clean image of it to the log, so recovery can always rebuild a torn page from a known-good base. This closes the single most common silent data-loss hole in naive storage engines.

- **Honest handling of disk-flush failures.** If the operating system ever reports that a flush failed, unidb refuses to keep claiming durability and requires a restart, rather than falsely reporting data as safe. After the restart, recovery still finds a consistent database.
- **Checksums on every page** catch bit-rot and partial writes on read.

3.5 Memory & space management

- **Buffer pool.** Hot pages are cached in a configurable in-memory pool (32 MiB by default). Working sets larger than the pool are handled safely — the engine flushes and evicts correctly rather than failing.
- **Free-space tracking.** Each table tracks its free space durably, so inserts stay fast and predictable even as tables grow into the millions of rows — insert cost stays flat where a naive design would degrade and eventually fail.
- **Vacuum & autovacuum.** Because updates create new row versions, superseded versions are reclaimed in the background once no transaction can still see them, keeping tables compact automatically (see Section 9).

4. Transactions & working with many users at once

4.1 Snapshots & isolation levels

unbdb uses **multi-version concurrency control (MVCC)**: an update never overwrites a row in place — it writes a new version and marks the old one as superseded. Each transaction reads from a consistent **snapshot**, so **readers never block writers and writers never block readers**. Old versions are cleaned up later, once no active transaction can still see them.

Three isolation levels are available:

Level	What you get	On conflict
Read Committed (default)	Each statement sees the latest committed data. The common, high-throughput default.	A write against a row another transaction is still changing is rejected so you can retry.
Repeatable Read	The whole transaction sees one stable snapshot — the same rows every time you read them.	Conflicting concurrent writes surface a serialization error; the caller retries.
Serializable	The strongest guarantee: the result is as if transactions ran one at a time, protecting against subtle write-skew anomalies.	A dangerous interleaving is detected and one transaction is asked to retry.

4.2 Reading and writing at the same time

unbdb is designed for real concurrency, and its behavior under load is straightforward to reason about:

- **Reads are lock-free.** Any number of readers proceed in parallel and never wait on each other or on writers.
- **Writers to different rows run in parallel** and scale across CPU cores. In benchmarks, raw write throughput scales **3.68× at 8 concurrent writers**, and concurrent inserts into indexed tables run about **25% faster** than the earlier serialized behavior.
- **Writers to the same row are safe by first-committer-wins.** If two transactions try to change the same row, one succeeds and the other gets a clear conflict error to retry — there is no lost update and no silent overwrite.
- **Deadlocks are detected automatically.** If two transactions wait on each other, the engine breaks the cycle by aborting one with a deadlock error, rather than hanging.
- **Per-query safety limits.** Queries can be given time budgets and can be cancelled, so one runaway query cannot monopolize the engine.

Under the hood, concurrent index updates are kept correct with fine-grained, carefully ordered locking so that inserts can proceed in parallel without ever producing an inconsistent index — validated by an extensive concurrency test suite (Section 11).

4.3 What happens on rollback

When a transaction is rolled back (explicitly, or because it hit an error or was abandoned), unbdb physically neutralizes everything it wrote: rows it inserted become invisible and rows it had marked superseded are restored. This is what lets graph edges, events, and large objects all inherit correct rollback behavior for free — they are ordinary rows, so undoing them uses the exact same machinery as undoing a table row.

5. The SQL layer

unbdb speaks a practical, PostgreSQL-flavored SQL dialect. Statements are parsed, planned, cost-optimized, and executed against versioned tables. You can inspect any plan with `EXPLAIN` (and its real row counts and timing with `EXPLAIN ANALYZE`).

5.1 Supported data types

SQL type	Stores	Notes
<code>INT</code> / <code>INTEGER</code> / <code>BIGINT</code>	64-bit signed integer	—
<code>SERIAL</code>	Auto-incrementing integer	Durable per-table counter; typical primary-key type.
<code>FLOAT</code> / <code>DOUBLE</code>	64-bit floating point	—
<code>DECIMAL(p,s)</code> / <code>NUMERIC(p,s)</code>	Exact fixed-point decimal	No precision loss — ideal for money.
<code>TEXT</code>	UTF-8 string	Can carry a full-text index.
<code>BOOL</code> / <code>BOOLEAN</code>	true / false	—
<code>JSON</code>	Nested JSON document	Returned as real nested JSON, not a string.
<code>UUID</code>	128-bit UUID	—
<code>DATE</code>	Calendar date	—
<code>TIME</code>	Time of day	—
<code>TIMESTAMP</code>	Date + time (UTC)	Serialized as a string so no precision is lost.
<code>BYTEA</code>	Raw binary bytes	—
<code>VECTOR(n)</code>	Fixed-length numeric vector	Embeddings for similarity search; can carry a vector index.

New types are added additively, without ever changing the on-disk format — existing databases keep working across engine upgrades.

5.2 Writing queries — examples

Create a table

```
CREATE TABLE users (  
  id          SERIAL PRIMARY KEY,  
  email       TEXT NOT NULL UNIQUE,  
  full_name   TEXT,  
  plan        TEXT DEFAULT 'free',  
  credits     DECIMAL(10,2) DEFAULT 0.00,  
  is_active   BOOL DEFAULT true,  
  created     TIMESTAMP,  
  CHECK (credits >= 0)  
);  
  
CREATE TABLE documents (  
  id          SERIAL PRIMARY KEY,  
  owner_id    INT REFERENCES users (id),  
  title       TEXT,  
  body        TEXT,  
  embedding   VECTOR(1536)  
);
```

Insert data (use bind parameters to stay injection-safe)

```
INSERT INTO users (email, full_name, plan)  
VALUES ($1, $2, $3);  
-- params: ["ada@example.com", "Ada Lovelace", "pro"]  
  
INSERT INTO documents (owner_id, title, body, embedding)  
VALUES (1, 'Notes', 'first draft', [0.12, 0.03, ...]);
```

Query — filters, joins, aggregates, sorting

```
-- Join + filter
SELECT u.email, d.title
FROM documents d
JOIN users u ON u.id = d.owner_id
WHERE u.plan = 'pro'
ORDER BY d.title
LIMIT 20;

-- Aggregate with grouping
SELECT plan, COUNT(*) AS n, AVG(credits) AS avg_credits
FROM users
WHERE is_active = true
GROUP BY plan
HAVING COUNT(*) > 5;

-- Subquery
SELECT email FROM users
WHERE id IN (SELECT owner_id FROM documents WHERE title LIKE 'Notes%');

-- Common table expression (CTE)
WITH active AS (SELECT id, email FROM users WHERE is_active = true)
SELECT * FROM active ORDER BY email;
```

Vector similarity search

```
-- Ten documents whose embedding is nearest to the query vector
SELECT id, title
FROM documents
WHERE NEAR(embedding, [0.10, 0.24, ...], 10);
```

`NEAR(column, query_vector, k)` returns the `k` nearest rows by vector distance and composes with ordinary `WHERE` filters and security policies.

Inspect a query plan

```
EXPLAIN ANALYZE
SELECT plan, COUNT(*) FROM users GROUP BY plan;
```

Traverse graph relationships (Cypher)

```
MATCH (a)-[:FOLLOWS]->(b)
WHERE a.id = 1
RETURN b.id;
```

5.3 Query capabilities

Area	Supported	Not yet supported
Joins	INNER / LEFT / RIGHT / CROSS, JOIN ... USING (cols) ; hash join (spills to disk on large inputs), sort-merge, and index-nested-loop	FULL OUTER, NATURAL

Aggregation & sort	COUNT / SUM / AVG / MIN / MAX, GROUP BY / HAVING, DISTINCT, ORDER BY (spills to disk), LIMIT / OFFSET	Window functions
Subqueries & CTEs	Scalar / IN / EXISTS (correlated and not), IN-list, non-recursive WITH	Recursive CTEs
Planning	ANALYZE (durable statistics: row counts, distinct counts, histograms), cost-based join ordering, index-vs-scan choice, EXPLAIN [ANALYZE]	—
Schema (DDL)	CREATE / ALTER / DROP / TRUNCATE, DROP COLUMN, SERIAL, prepared statements with \$n parameters	Views, triggers, stored procedures

Join, aggregate, and subquery results are continuously cross-checked against SQLite on shared data to guarantee correctness.

5.4 Constraints, security & roles

- **Constraints:** PRIMARY KEY , FOREIGN KEY / REFERENCES , UNIQUE , NOT NULL , CHECK , and DEFAULT — enforced on every insert and update.
- **Row-level security (RLS):** attach a policy predicate to a table (e.g. tenant_id = 7) and it is automatically applied to every query on that table — so multi-tenant isolation is enforced by the engine, not by every query author. It composes with vector search and index paths alike.
- **Users, roles & privileges:** create users and roles, grant per-table privileges (with role inheritance), and run statements as a given user so those privileges are enforced. Over the REST server, the caller's identity comes from their signed token.
- **Injection-safe by construction:** prepared statements with \$n bind parameters pass values as data, never as re-parsed SQL.


6. Search, indexing & relationships

6.1 Automatic durable indexes

Every secondary index in unidb — B-tree, vector, full-text, and graph adjacency — is **durable, crash-safe, and kept in step with your data inside the same transaction**. Indexes are never lost or rebuilt on open, so opening a database is instant no matter how large it is. Index-assisted queries always re-check results against your transaction's snapshot and security policy, so an index can only ever make a query *faster*, never return a wrong or unauthorized row.

6.2 Vector similarity search

A `VECTOR(n)` column can be given a durable similarity index, queried with the `NEAR` operator. unidb clusters vectors into cells, probes the most promising cells for a query, then re-ranks the candidates exactly — so there is **no approximation error in the final ranking**. In testing it returns the same top results as a brute-force search ($\text{recall@10} = 1.000$) while opening instantly and surviving crashes with its results intact.



Distance metrics: Euclidean (default) and cosine. The index is built once and kept durable — never recomputed on open.

Figure 7 — The `NEAR` vector-search path: probe promising cells, then re-rank exactly.

6.3 Full-text search

A `TEXT` column can carry a durable full-text (inverted) index. Search tokenizes your query and intersects the matching documents, then checks each candidate against your snapshot and security policy. Raw lookups run in the tens of microseconds.

6.4 Graph relationships

- **Edges are first-class but cost nothing extra to make safe.** A directed edge (`from, to, type, properties`) is stored as an ordinary row, so per-edge locking, rollback, and crash safety come for free.
- **Fast neighbor lookups.** A durable adjacency index finds a vertex's edges, and unidb batches the row fetches by page so a hot, highly-connected vertex costs one page read per ~128 edges instead of one per edge — about **9x faster** than the naive approach, bringing graph traversal to parity with Postgres.
- **Cypher reads:** `MATCH (a)-[:TYPE]->(b) WHERE ... RETURN ...` for single-hop, directed traversals. Edges are created and removed through the engine's API.

7. Real-time events & streaming

unidb has a durable event queue built in. When you opt a table into event capture, every insert, update, and delete on that table also durably records an event — **inside the same transaction as the change itself**. That

inline, transactional capture gives you properties an external message queue cannot:

- **Events can never disagree with your data.** An event is committed with the change that produced it; if that transaction rolls back, its events vanish with it. There is no dual-write gap.
- **No data loss across a crash.** Consumer offsets are stored durably and advanced only when a consumer acknowledges. A crash mid-delivery redelivers rather than dropping — at-least-once delivery is proven, not promised.
- **Kafka-style consumption:** poll from your durable offset, process, then acknowledge. A cleanup pass reclaims events once *every* registered consumer has moved past them — a slow consumer holds its events until it catches up, and cannot be starved.
- **Realtime fan-out.** The REST server exposes the stream as Server-Sent Events (live tail with resume-from-offset), and a companion dispatcher can push events to webhooks (with automatic retry and a dead-letter table) or broadcast rooms — the durable-by-default alternative to a typical realtime layer.

Measured, unidb's `poll_events` cycle runs in tens to hundreds of microseconds versus a few milliseconds for the equivalent Postgres `SKIP LOCKED` pattern (see Section 12 for the honest caveat on very large backlogs).

8. The REST server — full API reference

The optional REST server exposes the whole engine over HTTP with JWT authentication, TLS, realtime streaming, and metrics. It is a thin wrapper over the embedded engine: by default each request runs as one complete `begin → execute → commit` transaction, and many requests execute concurrently against one shared engine. The default embedded library has no network dependencies — the server is a separate, opt-in build.

8.1 Getting started

- **Base URL:** `http://127.0.0.1:8080` by default (configurable). Set a TLS certificate and key to serve HTTPS.
- **Authentication:** every route except `GET /metrics` requires a bearer token: `Authorization: Bearer <jwt>`. Tokens are verified (HS256) against a shared secret; the server does not issue them (plug in your own auth service). The token's `sub` claim is the acting username for per-user privileges.
- **Content type:** JSON in and out, except the raw-row routes (`/rows...`) which carry opaque bytes.
- **Errors:** every non-2xx JSON response has the shape `{ "error": "message", "code": "MACHINE_CODE" }` — see [§8.8](#).

8.2 Running SQL

POST `/sql`

Execute one or more `;`-separated SQL statements atomically. If any statement fails, the whole request rolls back.

Body

```
{
  "sql": "INSERT INTO users (email, plan) VALUES ($1, $2)",
  "params": ["ada@example.com", "pro"], // optional, injection-safe $n binds
  "isolation": "serializable",          // optional: read_committed | repeatable_read |
serializable
  "cursor": false                       // optional: true buffers a large result, returns a
cursor
}
```

Bind values: a JSON string binds as text (coerced to the column type), a number as int/float, a numeric array as a vector.

Response `200 OK` — one result object per statement, in order:

```
{
  "results": [
    { "type": "inserted", "count": 1 }
  ]
}
```

Result shapes include `created_table`, `created_index`, `inserted / updated / deleted / truncated` (with a count), `altered_table`, `dropped_table`, and for queries:

```
{
  "type": "rows",
  "columns": ["id", "name", "profile"],
  "rows": [ [1, "alice", { "status": "active" }] ]
}
```

A `JSON` column returns nested JSON; a `DECIMAL` returns a decimal string (e.g. `"9.90"`) and a `TIMESTAMP` a UTC string, so no precision is lost. The same route serves joins, aggregates, `GROUP BY / HAVING`, subqueries, `WITH`, `ANALYZE`, and `EXPLAIN [ANALYZE]` — no separate routes.

Cursor mode (`"cursor": true`, one rows-producing statement only) buffers the result server-side and returns `{ "cursor_id": 7, "columns": [...], "row_count": 120000 }` — page it with `GET /sql/cursor/{id}`.

GET `/sql/cursor/{cursor_id}` · **DELETE** `/sql/cursor/{cursor_id}`

Page (or drop) a cursor from a `cursor: true` request.

Query params: `limit` — rows per page (default 1000, max 10 000).

Response 200 OK:

```
{ "columns": ["id"], "rows": [[1],[2]], "done": false, "remaining": 118000 }
```

The final page reports `"done": true` and drops the cursor. Cursors are bound to the creating principal and expire after 60s idle; `DELETE` drops one early (204).

POST `/cypher`

Run a Cypher-subset graph read.

```
// Body
{ "query": "MATCH (a)-[:FOLLOWS]->(b) WHERE a.id = 1 RETURN b.id" }

// Response 200 OK
{ "results": [ { "type": "rows", "columns": ["id"], "rows": [[2],[3]] } ] }
```

8.3 Transactions over HTTP

For work spanning multiple requests, open a **transaction session**. Pass its id on later requests via the `X-Txn-Id` header; those requests run in the transaction and do **not** auto-commit.

POST `/txn/begin`

```
// Body (optional; empty = read_committed)
{ "isolation": "read_committed" | "repeatable_read" | "serializable" }

// Response 201 Created
{ "txn_id": 42, "isolation": "read_committed",
  "idle_timeout_secs": 60, "expires_at": "2026-07-11 12:34:56" }
```

`expires_at` is a sliding idle deadline — each request on the session pushes it out again.

POST `/txn/{txn_id}/commit` · **POST** `/txn/{txn_id}/rollback`

```
// Response 200 OK
{ "txn_id": 42, "state": "committed" } // or "rolled_back"
```

Session rules: one request at a time per session (a second concurrent request gets `409 TXN_BUSY`); a session belongs to the principal that created it (`403 TXN_FORBIDDEN` otherwise); abandoned sessions are auto-aborted after an idle timeout (default 60s) and return `404 TXN_NOT_FOUND`; DDL is not allowed inside a session (`400 DDL_IN_SESSION` — run it as a one-shot request); a failed mutating statement rolls the session back. Sessions do not survive a server restart.

8.4 Rows, graph & indexes

The `/rows...` routes store **opaque bytes** (raw payloads unidb does not interpret) — use `/sql` for typed, columnar data. All accept `X-Txn-Id` to join a session.

Route	Purpose & payload	Response
POST <code>/rows</code>	Insert one raw row; body is raw bytes.	201 { "row_id": { "page_id": 3, "slot": 0 } }
POST <code>/rows/batch</code>	Insert up to 10 000 raw rows atomically; body { "rows": ["<base64>", ...] } (32 MiB total).	201 { "row_ids": [...] }
GET <code>/rows/{page_id}/{slot}</code>	Read a row by id.	200 raw bytes, or 404 NOT_FOUND
PUT <code>/rows/{page_id}/{slot}</code>	Update a row; body is the new raw bytes.	200 { "row_id": { ... } } (id may change)
DELETE <code>/rows/{page_id}/{slot}</code>	Delete a row.	204 No Content
POST <code>/edges</code>	Create an edge; body { "from_id": 1, "to_id": 2, "edge_type": "FOLLOWS", "props": { ... } } (props optional).	201 { "row_id": { ... } }
DELETE <code>/edges/{page_id}/{slot}</code>	Delete an edge; body { "from_id": 1 } (required).	204 No Content
GET <code>/edges/from/{from_id}</code>	List outgoing edges from a vertex.	200 { "edges": [{ "row_id", "to_id", "edge_type", "props" }] }
POST <code>/indexes</code>	Create an index; body { "table": "documents", "column": "embedding", "kind": "Hnsw" } (kind: "Hnsw" on a vector column or "FullText" on a text column; omit kind to drop).	204 No Content
GET <code>/indexes/{table}/{column}/status</code>	Report a column's index status.	200 { "status": "Ready" } or { "status": null }

GET /tables	List user tables and their columns (cheap, no scan). For full introspection query <code>information_schema.* over /sql</code> .	200 array of { "name", "columns": [{ "name", "type", "nullable", "index" }] }
--------------------	---	---

8.5 Events & realtime

Route	Purpose & payload	Response
POST /tables/{table}/events	Opt a table into event capture. No body.	204 No Content
POST /events/ack	Advance a consumer's durable offset; body { "consumer": "billing-worker", "up_to_seq": 17 }.	204 No Content
POST /events/vacuum	Reclaim events acknowledged by every consumer. No body.	200 { "reclaimed": 17 }

GET /events/subscribe (Server-Sent Events)

Stream new events as SSE frames. Two modes:

- **Durable consumer** (`consumer` set): at-least-once, resumes from that consumer's acked offset; un-acked events are re-yielded until acked.
- **Ephemeral live-tail** (`consumer` omitted): browser-style tail. Resumes past the standard `Last-Event-ID` reconnect header, else `?from_seq=`, else the start.

Query params: `consumer`, `from_seq`, `table` (filter), `limit` (per poll, default 100), `interval_ms` (poll interval, default 500).

Response 200 OK, `text/event-stream`:

```
id: 17
event: insert
data: {"seq":17,"xid":42,"table_name":"orders","op":"insert","payload":{"id":1,"total":9.99}}
```

Acknowledge separately via `POST /events/ack` — acks are not sent over the stream.

8.6 Admin, stats, logs & metrics

Route	Purpose & payload	Response
PUT /tables/{table}/rls	Attach a row-level-security policy (superuser); body { "predicate": "tenant_id = 7" } (an AND-only WHERE subset).	204 ; 400 SQL_PARSE_ERROR / 403 PERMISSION_DENIED
POST /admin/flush	Force the log durable, then flush all dirty pages (superuser). No body.	204 No Content
POST /checkpoint	Checkpoint: flush, write a checkpoint record, trim the log. No body.	204 No Content
GET /stats	Activity + health snapshot (see below).	200 JSON

GET /logs	Bounded, cursor-paged tail of structured JSON logs (superuser). Params: level, since, until, q, cursor, limit (≤500).	200 { "logs": [...], "returned", "scanned", "truncated", "next_cursor" }
GET /metrics	Prometheus exposition. The only route with no auth.	200 text/plain

GET /stats is a pg_stat_*-style snapshot: commit / abort / checkpoint counts, active transactions, WAL bytes, per-statement-kind latency percentiles, WAL-flush latency, buffer-pool hit ratio, lock waits and deadlocks, the vacuum-horizon age, per-table page counts, parallel-worker utilization, autovacuum activity, recent slow queries, and open sessions / cursors. Example fragment:

```
{
  "commits": 42, "aborts": 3, "active_transactions": 0, "wal_bytes": 81920,
  "statement_latency": { "insert": { "count": 50, "p50_us": 32, "p99_us": 256 } },
  "bufferpool": { "hits": 980, "misses": 40, "hit_ratio": 0.9607 },
  "locks": { "waits": 0, "deadlocks": 0 },
  "horizon_age_secs": 0.0,
  "tables": [ { "name": "t", "pages": 3 } ]
}
```

8.7 Replication

Route	Purpose
POST /replication/slots	Create a slot; body { "name": "...", "sync": false } → 201 { "name", "restart_lsn", "kind" } .
GET /replication/slots	List slots → { "slots": [...] } .
DELETE /replication/slots/{name}	Drop a slot → 204 .
POST /replication/slots/{name}/advance	Confirm applied up to an LSN; body { "lsn": <n> } → 204 .
GET /replication/stream?from_lsn={n}	Ship log records after from_lsn as a byte stream; the primary's tail LSN is in the x-unidb-tail-lsn response header.

8.8 Error codes

Every error carries a stable machine-readable code :

HTTP	Code	Meaning
404	TABLE_NOT_FOUND / COLUMN_NOT_FOUND / NOT_FOUND	Referenced table/column doesn't exist; row has no visible version
404	TXN_NOT_FOUND / CURSOR_NOT_FOUND	Unknown/finished/expired session or cursor id
409	TABLE_ALREADY_EXISTS	CREATE TABLE on an existing name
409	WRITE_CONFLICT / SERIALIZATION_FAILURE / DEADLOCK	Concurrency conflict — retry the transaction
409	UNIQUE_VIOLATION	Duplicate UNIQUE / PRIMARY KEY value
409	TXN_BUSY	Second concurrent request on one session
408	QUERY_TIMEOUT / QUERY_CANCELLED	Per-query time budget or cancellation

403	PERMISSION_DENIED / TXN_FORBIDDEN / CURSOR_FORBIDDEN	Missing privilege / not your session or cursor
400	SQL_PARSE_ERROR / SQL_PLAN_ERROR / SQL_UNSUPPORTED	Malformed or unsupported SQL
400	NOT_NULL_VIOLATION / CHECK_VIOLATION / FOREIGN_KEY_VIOLATION	Constraint failed
400	DDL_IN_SESSION / ISOLATION_IN_SESSION / CURSOR_NOT_ROWS	Operation not allowed in this context
400	EMPTY_BATCH / BAD_BASE64 / BATCH_T00_LARGE	Invalid batch-insert payload
401	UNAUTHORIZED	Missing/invalid/expired token
503	DURABILITY_FAILURE	A disk flush failed; the engine must be restarted (durability can no longer be guaranteed)
500	INTERNAL_ERROR	I/O, checksum, or corruption a well-formed request should never trigger

9. Operations & high availability

- **Segmented write-ahead log.** The log is kept as fixed-size segments; consumed segments are deleted whole. Checkpoints run automatically on a time or size trigger (at a quiescent point) to keep the log bounded and recovery fast.
- **Replication & read replicas.** A replica seeds from a base snapshot and applies shipped log incrementally. The primary ships only data it has made durable, so a replica is always a consistent prefix of the primary — it can never show a change the primary might lose. On failover, a replica is promoted to read-write, with an optional synchronous-commit gate for zero-loss failover.
- **Backups & point-in-time restore.** Take a base backup plus archived log and restore to a chosen point. Measured on a sample database: base backup 7 ms, restore 72 ms, failover 26 ms.
- **Autovacuum.** A background worker reclaims dead row versions once no reader (including replicas) can still see them, scrubbing every index before space is reused. Under heavy churn it keeps a table compact automatically (bounded at ~35 pages vs ~82 unvacuumed in one test) and keeps point reads fast.
- **Security & observability.** Users, roles and privileges; verified JWT identities; native TLS; an append-only audit log; lock-free health metrics over `/stats` and `/metrics`; and structured JSON logs with a request id that correlates a single request across the app, slow-query, and audit logs (ready to ship to CloudWatch / Datadog / Loki). An operations runbook ships with the engine.

10. Performance — measured against Postgres

All numbers below are from release builds on Apple Silicon at real flush-to-disk durability, measured with the project's benchmark harness against Postgres (with pgvector, at matched durability) and, for the headline, against the full four-system stack unidb replaces. unidb reports honest wins *and* the places it still trails.

10.1 The multi-model advantage, quantified

Measurement	Result
Save a row + its embedding + a graph edge + an event, one durable commit, vs. the same across Postgres + a vector store + a graph table + an outbox (four commits)	3.61× faster (250 vs 69 txns/s) at real durability
Crash consistency across those four writes	Unconditional win: 0 orphans vs. a torn record
Adding each extra model to the same commit	Four models cost about the same as one (they share a single flush)

10.2 Where unidb wins on single-model work

Workload	Result vs. Postgres
<code>SELECT COUNT(*)</code>	2.81× faster
Filtered scans (parallel workers, on by default)	6.4–6.6× faster ; full scan up to ~35.7M rows/s at 1M rows
Event polling (<code>poll_events</code> vs <code>SKIP LOCKED</code>)	~20–980 µs vs ~2.6–3.1 ms
Embedded point reads	~5× faster (~1.09 µs)
Graph adjacency (hot vertex)	Parity (~94 µs vs ~98 µs at 1k edges)

10.3 What made it fast

Improvement	What it does	Measured effect
Group commit	Many transactions share one durability flush; read-only commits skip it	Multi-model commit 33.1 → 4.3 ms (~7.7×); server 135 → 4,780 requests/s at 50 clients; point read 3.05 ms → ~1.09 µs
Concurrent writers	Thread-safe engine, fine-grained latches, real lock manager	Raw writes scale 3.68× at 8 writers; indexed inserts +25%
Parallel scans (default on)	Scan work partitioned across CPU cores over the shared memory map	COUNT 3.82×, filtered scans up to 6.6×
Partial-decode reads	Decode only the columns a query needs; count from headers alone	<code>COUNT(*)</code> 2.81× faster than Postgres
Durable free-space map	Per-table space tracking; instant open map	Insert cost flat ~17–28 µs/row into the millions (was degrading, then failing)
Batched graph resolution	Fetch a vertex's edges grouped by page (~128 edges/page)	~9× over naive; parity with Postgres
Durable vector index	Similarity index kept on disk, never recomputed on open	recall@10 = 1.000; open instant (vs 30 s in-memory rebuild)

10.4 Scale

Bulk insert holds ~31k rows/s flat to 2M rows; a full-table scan runs at ~17.5M rows/s at 1–2M rows; the engine handles databases of ten million rows and beyond. The confirmed target is a strong single node plus read replicas — hundreds of GB with high read throughput.

11. Correctness & testing

unbdb's safety claims are backed by tests, not assertions of good intent:

- **Crash testing.** A dedicated harness kills the engine at dozens of precise points — after a log write but before a flush, mid-checkpoint, mid-recovery, during log trimming, mid-vacuum, on a torn page, on a simulated disk-flush failure, and after total data-file loss — then reopens and asserts the database contains exactly the committed transactions. A property test drives random transaction sequences with random crash points under both durability settings.
- **Concurrency testing.** A production-shaped concurrency matrix checks exact visible-row sets, no duplicate rows in any snapshot, agreement between index and scan paths, and repeatable reads — under CPU contention and with hang deadlines.
- **Differential testing.** Join, aggregate, and subquery results are compared against SQLite on shared data.
- **Every change is gated.** Formatting, linting, the full test suite, the crash harness, and recorded benchmarks must all pass before any change lands.

The engine is rigorously tested and crash-safe by design, but it is still early-stage software — it has not accumulated the decades of production hardening that Postgres has. Section 12 is an honest account of where it stands.

12. Known limitations, bugs & roadmap priority

unidb keeps an honest, public account of its gaps — both missing features and known bugs — so nothing below is a surprise in production. Every item carries a **priority** (how much it matters to close) and a status: actively **Tracked**, or explicitly **Planned** for a future phase.

12.1 Known bugs (open, tracked)

This is the complete list of currently open, user-visible correctness or behavior bugs. It is short by design: every other correctness bug uncovered during development — including a subtle transaction-rollback ordering issue found under heavy concurrent load — was root-caused, fixed, and is now locked in place by a permanent automated regression test (Section 11). Nothing below causes silent data loss or a wrong answer; each is a query that is rejected when it should be accepted.

Bug	What happens	Priority
Inconsistent IS NULL / IS NOT NULL filter support	WHERE column IS NULL works on a query that already sorts, joins, or aggregates, but the identical filter is rejected as unsupported on a plain filtered lookup — and, notably, when combined with a vector similarity search (NEAR(...) AND column IS NULL). The same predicate silently behaves differently depending on unrelated parts of the query, which is a genuine surprise for anyone hitting it.	High — blocks a common "similarity search + filter" pattern

12.2 Known limitations & feature gaps

Area	Limitation	Priority	Status
Transactional DDL	A CREATE TABLE in a transaction that later aborts is not rolled back through crash recovery (per-request rollback does apply today)	High	Tracked
Serializable isolation	Row-granularity only — no predicate locks, so no phantom protection; sound but occasionally over-conservative	High	Tracked
Auth platform	The server verifies tokens but does not issue them — no built-in sign-up/login/OAuth	High	Planned — next phase
SQL surface	No window functions, recursive CTEs, FULL OUTER / NATURAL joins, views, triggers, or stored procedures; single-column indexes only	Medium	Planned — next phase
Event queue at scale	Polling cost grows with table size (no sequence index yet); realtime is poll-per-subscriber rather than push	Medium	Planned — next phase
Vacuum	Engine-global (not per-table) accounting; no cost-based throttle; no whole-table compaction	Medium	Tracked
Replication	Point-in-time restore is by log position, not wall-clock time; no logical replication yet	Medium	Tracked
Updates	An UPDATE writes a new row version (more write volume than Postgres's in-place update path); measured ~0.34× of Postgres on bulk updates	Low	Planned — next phase
Encryption at rest	Not yet available	Low	Planned — next phase
Graph	Cypher is single-hop, read-only; no reverse-direction traversal yet	Low	Tracked

Maturity	Rigorous and crash-tested, but early-stage — not yet battle-hardened over decades like Postgres	Low	Ongoing
----------	---	-----	---------

13. Where unidb is headed

The roadmap closes the gaps that block production adoption while protecting the one thing that makes unidb different — multi-model atomicity. It is framed against the two systems users compare unidb to: **Postgres** (as an engine) and **Supabase** (as a backend platform).

13.1 Toward Postgres-level engine completeness

Theme	What's coming
Transactional completeness	Transaction-scoped, crash-safe DDL; savepoints and nested rollback; full serializable isolation with phantom protection; complete foreign-key actions (<code>ON DELETE / UPDATE CASCADE</code> , etc.)
SQL surface	Views and materialized views; window functions; triggers and stored procedures; recursive CTEs and remaining join types; richer types (arrays, ENUM, INTERVAL, ranges); multi-column, partial, and expression indexes
Performance parity	In-place (HOT-style) updates to close the update gap; more query work pushed onto parallel workers; a sequence index for the event queue
Ecosystem & operations	A PostgreSQL wire-protocol listener so existing drivers, ORMs, and BI tools connect unchanged; time-based point-in-time restore; logical replication / change-data-capture connectors; encryption at rest

13.2 Toward a backend-in-a-box (Supabase-style platform)

Several of these have already shipped; others are planned.

Capability	Today	Direction
Realtime	Durable SSE stream with resume + webhook / broadcast fan-out	Shipped — durable-by-default, a genuine advantage over realtime layers that aren't
Observability	Lock-free per-chokepoint metrics over <code>/stats</code> and <code>/metrics</code>	Shipped
Logs	Structured JSON logs with request-id correlation and a bounded <code>/logs</code> tail	Shipped
Auth platform	Token verification, users, roles, per-table grants	Planned — a full auth service (sign-up / login, OAuth, magic links, MFA) in front of the existing core
Client SDKs	REST + Rust attach client	Planned — TypeScript SDK first, then Python, generated from the REST contract
Auto REST richness	Hand-built routes; SQL over HTTP	Planned — PostgREST-style filtering / ordering / embedding, OpenAPI generation
Object storage	Transactional large objects inside the engine (multi-GB), tiered to S3/MinIO-compatible storage with presigned upload/download and automatic reconciliation	Shipped
Dashboard / studio	—	Planned — a web console: schema browser, SQL editor with EXPLAIN, metrics, RLS/role editor

Migrations tooling	SQL DDL + docs	Planned — versioned migration files and a CLI once transactional DDL lands
Edge functions / hosting	—	Not planned — no cloud control plane by design; run unidb embedded inside your own function runtime

13.3 Deliberately parked

- **Columnar / analytical engine** — the opposite axis to the multi-model transactional thesis; gated on real analytics demand.
- **Distributed / sharded multi-writer** — reverses the single-primary design and strains cross-model atomicity; the confirmed target remains a strong single node plus read replicas.

unidb — Engine Architecture & Product Guide. This is a distilled reference; where it disagrees with the engine's canonical source docs, those win. Editing and regeneration instructions (and the source-material provenance) are in `unidb_engine_architecture_context.md`, alongside this file.